

# MIT6.828 Lab2: Memory Management

Zhuofan Zhang

Jan 2022

# Contents

- 1 Physical Page Management 1
- 2 Virtual Memory 6
  - 2.1 Virtual, Linear, and Physical Addresses . . . . . 6

# Chapter 1

## Physical Page Management

本次 Lab 的内容主要分为两个部分，第一个部分是对可用物理内存的管理，第二个部分是虚拟内存的管理。

操作系统负责对可用物理内存的管理。JOS 的物理页使用 `PageInfo` 结构相关数据（数组/空闲链表）进行物理内存管理。`PageInfo` 数组（`pages`）是对整个物理地址空间（32 位，4GB）的映射：将整个物理地址空间按 4KB-page 划分，则每个单页可由 `pages` 对应下标的 `PageInfo` 结构唯一指代（原理可见 `pmap.h/page2pa` 的实现）。

### Exercise 1

In the file `kern/pmap.c`, you must implement code for the following functions:

**`boot_alloc()`**

**`mem_init()` (only up to the call to `check_page_free_list(1)`)**

**`page_init()`**

**`page_alloc()`**

**`page_free()`**

`check_page_free_list()` and `check_page_alloc()` test your physical page allocator. You should boot JOS and see whether `check_page_alloc()` reports success. Fix your code so that it passes. You may find it helpful to add your own `assert()`s to verify that your assumptions are correct.

根据题目要求，我们依次实现上述函数。

对于 **`boot_alloc`** 函数，根据注释提示可知其为系统启动初期的 `physical-page allocator`，协助完成初始化工作。函数中使用了一个变量 `end`，我们从内核的链接脚本（`kernel.ld`）中可以定位到它的位置，即 `.BSS` 段的末尾；根据注释提示，`end` 表示第一处未被 `kernel` 使用的虚拟地址。由 Lab1 可知内核被链接至高地址位（`0xf0100000` 开始的 4MB 区域），因此也许保证在这个阶段分配的物理内存不超过 4MB，需要对分配是否越界进行检查。明确上述要求后，我们即可实现该函数：

```

static void *
boot_alloc(uint32_t n)
{
    static char *nextfree; // virtual address of next byte of free memory
    char *result;

    // Initialize nextfree if this is the first time.
    // 'end' is a magic symbol automatically generated by the linker,
    // which points to the end of the kernel's bss segment:
    // the first virtual address that the linker did not assign
    // to any kernel code or global variables.
    if (!nextfree) {
        extern char end[];
        nextfree = ROUNDUP((char *) end, PGSIZE);
    }

    // Allocate a chunk large enough to hold 'n' bytes, then update
    // nextfree. Make sure nextfree is kept aligned
    // to a multiple of PGSIZE.
    //
    // LAB 2: Your code here.
    result = nextfree;
    if(n > 0)
        nextfree = ROUNDUP(nextfree + n, PGSIZE);
    else if(n < 0)
        // n < 0: should not come here
        panic("boot_alloc: parameter n < 0.");

    // fix me: need to be explained
    if(PADDR(nextfree) >= 0x400000)
        panic("boot_alloc: run out of mem.");

    return result;
}

```

**mem\_init** 函数是被内核 `i386_init()` 调用的函数，用来完成内核内存管理的初始化；本节仅需完成物理内存管理相关的部分。

根据本节开头我们提到的物理内存管理方法，我们首先需要创建管理物理页的数组 `pages`。由于 `pages` 本身也需要占用内存，我们利用上一步实现的 `boot_alloc` 分配内存并将 `pages` 进行 0 初始化。需要注意的是数组大小 (`npages`) 由 `i386_detect_memory` 调用得到，它负责检查系统可用的物理内存大小。根据上述分析，我们实现 `mem_init` 中初始化 `pages` 部分的内容：

```

void
mem_init(void)
{
    ...

    // Find out how much memory the machine has (npages & npages_basemem).
    i386_detect_memory();

    ...

    //////////////////////////////////////
    // Allocate an array of npages 'struct PageInfo's and store it in 'pages'.
    // The kernel uses this array to keep track of physical pages: for
    // each physical page, there is a corresponding struct PageInfo in this
    // array. 'npages' is the number of physical pages in memory. Use memset
    // to initialize all fields of each struct PageInfo to 0.
    // Your code goes here:
    pages = (struct PageInfo *) boot_alloc(npages * sizeof(struct PageInfo));
    memset(pages, 0, npages * sizeof(struct PageInfo));

    ...
}

```

上一步实现的代码中，我们可以看到在为 `pages` 分配内存并进行 0 初始化后，`mem_init` 调用了真正对 `pages` 及空闲列表 `page_free_list` 进行初始化的 `page_init`。由于实际上系统运行至这一步时已经使用了一部分物理页，因此我们必须根据系统当前的状态正确设置页内容，并将空闲页放入空闲列表。根据 `page_init` 中的提示，我们完成代码如下。其中比较需要关注的是 `extended memory` 空闲部分的初始化，第三步除了设置 IO hole 不空闲外，也需设置已经被使用的物理页，而确定当前已用物理页位置末端（以虚拟地址表示）的方法是对 `boot_alloc` 的零调用：

（注：该部分内容可以结合 Lab1 中给出的 FIGURE: The PC's Physical Address Space 理解）

```

void
page_init(void)
{
    // The example code here marks all physical pages as free.
    // However this is not truly the case. What memory is free?
    // 1) Mark physical page 0 as in use.
    //     This way we preserve the real-mode IDT and BIOS structures
    //     in case we ever need them. (Currently we don't, but...)
    // 2) The rest of base memory, [PGSIZE, npages_basemem * PGSIZE)
    //     is free.
    // 3) Then comes the IO hole [IOPHYSMEM, EXTPHYSMEM), which must
    //     never be allocated.
    // 4) Then extended memory [EXTPHYSMEM, ...).
    //     Some of it is in use, some is free. Where is the kernel
    //     in physical memory? Which pages are already in use for
    //     page tables and other data structures?
    //
    // Change the code to reflect this.
    // NB: DO NOT actually touch the physical memory corresponding to
    // free pages!
    // 1) Mark physical page 0 as in use
    pages[0].pp_ref = 1;
    // 2) Base-memory
    size_t i;
    for (i = 1; i < npages_basemem; i++) {
        pages[i].pp_ref = 0;
        pages[i].pp_link = page_free_list;
        page_free_list = &pages[i];
    }
    // 3) IO hole + pages that have been used
    int hole_and_used = npages_basemem +
        ((EXTPHYSMEM - IOPHYSMEM) +
        PADDR(boot_alloc(0)))/PGSIZE;
    for(; i < hole_and_used; i++)
        pages[i].pp_ref = 1;

    // 4) extended memory
    for(; i < npages; i++)
    {
        pages[i].pp_ref = 0;
        pages[i].pp_link = page_free_list;
        page_free_list = &pages[i];
    }
}

```

实现了 `pages` 及 `page_free_list` 的初始化后,就可以进一步实现分配和回收页的函数:`page_alloc` 与 `page_free`。注意到实现物理页分配时,对页面进行初始化时 需要获得物理页的虚拟地址 (C语言中地址均匀虚拟地址),根据提示使用 `page2kva` 实现。`page2kva` 本身的实现是使用 `page2pa` 得到物理页的实际地址,再根据当前的映射使用 `KADDR` 转化为虚拟地址。

```
struct PageInfo *
page_alloc(int alloc_flags)
{
    // Fill this function in
    if(!page_free_list)
        return NULL;

    struct PageInfo *alloc_page = page_free_list;
    page_free_list = page_free_list->pp_link;
    alloc_page->pp_link = NULL;

    if(alloc_flags & ALLOC_ZERO)
        memset(page2kva(alloc_page), 0, PGSIZE);

    return alloc_page;
}

void
page_free(struct PageInfo *pp)
{
    // Fill this function in
    // Hint: You may want to panic if pp->pp_ref is nonzero or
    // pp->pp_link is not NULL.
    if(pp->pp_ref != 0 || pp->pp_link != NULL)
        panic("page_free: pp_ref != 0");

    pp->pp_link = page_free_list;
    page_free_list = pp;
}
```

# Chapter 2

## Virtual Memory

### Exercise 2

Look at chapters 5 and 6 of the Intel 80386 Reference Manual, if you haven't done so already. Read the sections about page translation and page-based protection closely (5.2 and 6.4). We recommend that you also skim the sections about segmentation; while JOS uses the paging hardware for virtual memory and protection, segment translation and segment-based protection cannot be disabled on the x86, so you will need a basic understanding of it.

这一节内容需要我们构建虚拟内存管理的基本设施，Exercise2 为我们提供了 80386 处理器内存管理模式的相关材料。阅读相关材料后，我们可以整理出处理器为我们提供的内存管理机制：分段/分页。

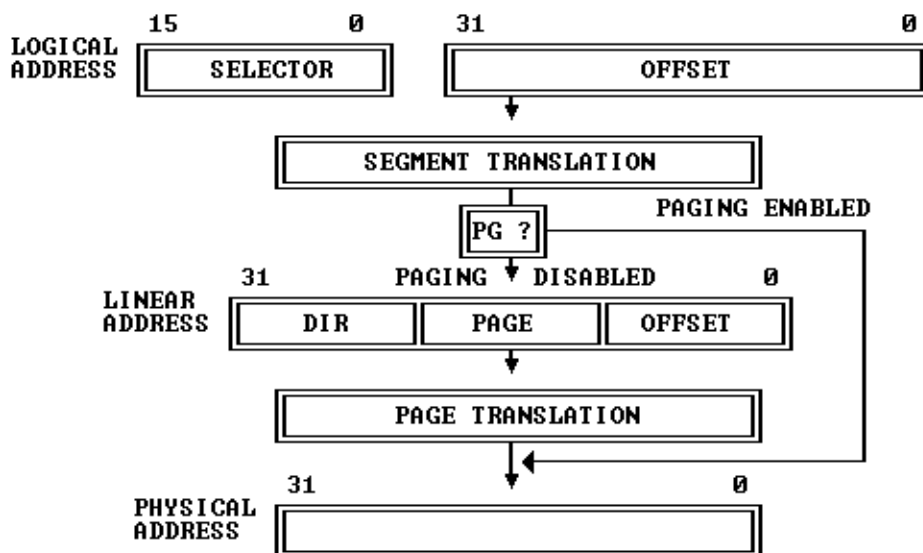
### 2.1 Virtual, Linear, and Physical Addresses

在 x86 体系中，我们可以定义 3 种地址：虚拟地址（Virtual Address, VA）[也称 Logical Address]，线性地址（Linear Address, LA）与物理地址（Physical Address, PA）。

对于用户而言，可见的只有 VA，而数据在物理内存中的实际地址是 PA。x86 的典型地址翻译流程如下图所示，其中  $VA \rightarrow LA$  的过程为分段（*Segmentation*）机制， $LA \rightarrow PA$  过程为分页（*Paging*）机制。

关于分页与分段的详细内容不在此处展开，但需知道：JOS 并不使用分段机制，但 x86 系统并没有显式关闭分段的方法。因此，在设置 Segment Descript Table 时，JOS 将段基址设置为 0，并将 LIMIT 设置为 0xffffffff，即将整个物理地址空间视为一段，在这种情况下  $VA = LA$ 。设置的过程在 Lab1 中的 boot 阶段完成（boot.S）。





### Exercise 3

While GDB can only access QEMU's memory by virtual address, it's often useful to be able to inspect physical memory while setting up virtual memory. Review the QEMU monitor commands from the lab tools guide, especially the `xp` command, which lets you inspect physical memory. To access the QEMU monitor, press `Ctrl-a c` in the terminal (the same binding returns to the serial console).

Use the `xp` command in the QEMU monitor and the `x` command in GDB to inspect memory at corresponding physical and virtual addresses and make sure you see the same data.

Our patched version of QEMU provides an `info pg` command that may also prove useful: it shows a compact but detailed representation of the current page tables, including all mapped memory ranges, permissions, and flags. Stock QEMU also provides an `info mem` command that shows an overview of which ranges of virtual addresses are mapped and with what permissions.

这一个练习的第一个要求与 Lab1 类似，要求我们观察虚拟地址与其实映射到的物理地址的内容是否一致（即虚拟内存映射是否成功）。与 Lab1 不同，此处采用更泛用的方法：使用 GDB 观察 PA，使用 QEMU-monitor 观察 VA。根据提示，我们在 gdb 设置内存映射成功后的位置检查，可以看到 VA 及其对应的 PA 内容的一致性：

```

zhangzf@DESKTOP-HPF7G3K:~/6.828/lab$ make qemu-gdb
***
*** Now run 'make gdb'.
***
/home/6.828/qemu/bin/qemu-system-i386 -drive file=obj/kern/kernel.img,index=0,media=disk,format=raw -serial mon:stdio -gdb tcp::26000 -D qemu.log -S
VMX server running on '127.0.0.1:5900'
QEMU 2.3.0 monitor - type 'help' for more information
(qemu) xp/4x 0xf00000
0000000000100000: 0x1badb002 0x00000000 0xe4524ffe 0x7205c766
(qemu)

zhangzf@DESKTOP-HPF7G3K:~/6.828/lab$ make gdb
gdb -n -x .gdbinit
GNU gdb (Ubuntu 9.2-0ubuntu1-20.04) 9.2
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB. Attempting to continue with the default i386 settings.

The target architecture is assumed to be i386
[f000:ffff] 0xffff0: jmp $0xf000,$0xe05b
0x0000ffff in ?? ()
+ symbol-file obj/kern/kernel
(gdb) b i386_init
Breakpoint 1 at 0xf0100040: file kern/init.c, line 15.
(gdb) c
Continuing.
The target architecture is assumed to be i386
-> 0xf0100040 <i386_init>: endbr32

Breakpoint 1, i386_init () at kern/init.c:15
15      if
(gdb) x/4x 0xf0100000
0xf0100000 <_start-260435460>: 0x1badb002 0x00000000 0xe4524ffe 0x7205c766
(gdb)

```

此外，我们还根据提示，使用 **info pg/mem** 命令查看页表情况与内存权限情况：

```

(qemu) info pg
VPN range      Entry      Flags      Physical page
[00000-003ff] PDE[000]    ----A----P
    [00000-000ff] PTE[000-0ff] -----WP 00000-000ff
    [00100-00100] PTE[100]     ----A---WP 00100
    [00101-00111] PTE[101-111] -----WP 00101-00111
    [00112-00112] PTE[112]     ---DA---WP 00112
    [00113-003ff] PTE[113-3ff] -----WP 00113-003ff
[f0000-f03ff] PDE[3c0]    ----A---WP
    [f0000-f00ff] PTE[000-0ff] -----WP 00000-000ff
    [f0100-f0100] PTE[100]     ----A---WP 00100
    [f0101-f0111] PTE[101-111] -----WP 00101-00111
    [f0112-f0112] PTE[112]     ---DA---WP 00112
    [f0113-f03ff] PTE[113-3ff] -----WP 00113-003ff
(qemu) info mem
0000000000000000-0000000000400000 0000000000400000 -r-
00000000f0000000-00000000f0400000 0000000000400000 -rw

```

本节的最后，Lab 还简要介绍了 JOS 中地址类型的 typedef。JOS 将地址类型区分为虚拟地址 (`uintptr_t`) 与物理地址 (`physaddr_t`)，它们均为 `uint32_t` 的 typedef。由于 C 中所有出现的“地址”均为虚拟地址，故其中 `uintptr_t` 可以被解引用使用，而 `physaddr_t` 不应被解引用，它仅作为物理地址字面值被使用，对其解引用将引发 UB。